

A Low-code approach for enforcing Fine-Grained Access Control in Database-centric applications

Final Project of Formal Model Driven Engineering

Student: Hoàng Nguyễn Phước Bảo

Universidad Autónoma de Madrid, Spain

Abstract. In this report, I propose a low-code approach to realize a model-driven proposal to enforce fine-grained access control in database-centric applications. The works involve designing the related metamodels and their domain specific languages, then, applying some model transformation and model manipulation techniques to generate final executable artifacts on the database level.

1 Introduction

Model-Driven Engineering (MDE) is a software development methodology that focuses on creating *models* of different views of a system, and then automatically generating different system artifacts from these models, such as code and configuration data. *Model-Driven Security* (MDS) is a specialization of model-driven engineering for developing secure systems. In a nutshell, designers specify system models along with their security requirements and use tools to automatically generate security-related system artifacts, such as access control infrastructures.

SQL Security Injector (SQLSI) is my first and latest work in the MDS discipline. The project attempts to enforce *fine-grained access control* (FGAC) in database using model-driven methodology. So far, this involves (i) to propose a formal model to express such access control policy and a rigorous design to enforce it in database-centric application[], then (ii) to realize the aforementioned work in a notably case study[], finally, (iii) to prove its correctness. In regards to the second goal, since the inputs in [] are indeed models of different kinds, I believe it makes sense to apply a more *low-code* approach, taking advantage of the knowledge I have learnt from the course, to replace the current *ad-hoc* realization represented in [].

Organization The rest of the report is organized as follows. In Section 2, I provide some main remarks about the SQLSI idea and design, including some related works and their drawbacks, from which uplift the approach in this report. Next, in Section 3, I describe the input and output of this new approach. Then, in Section 4, provides the in more detail the realization and technology used. Finally, in Section 5, I provide conclusion and possible future works.

2 Background and Motivation

2.1 Traditional approach vs SQLSI approach

Figure 1 shows the main difference between traditional and model-driven approach in enforcing FGAC in database-centric application. In particular:

Traditional approach On the left-hand side is the traditional approach to enforce FGAC in database-centric applications. Firstly, the policy will be manually implemented and enforced on the application level. Next, given a SQL statement, the application layer sends it to the data layer from which returns the result. Then, the application layer must enforce the FGAC policy to the result and finally return the answer to the client.

SQLSI approach On the right-hand side is the SQLSI model-based approach. This approach requires an application *data-model* that corresponds to the underlying database in the data layer and a *security-model* that defines the FGAC policy and respects the given data-model. The SQLSI component have two main functionalities:

- (i) Given those data- and security-models, SQLSI component generates a set of SQL *authorization functions* which are executable in the database.
- (ii) Then, instead of issuing directly the SQL statements from the application layer, SQLSI will *rewrite* these statements into semantically-equivalent SQL *secure stored-procedures* in which *injects* the given access control policy via means of aforementioned SQL *authorization functions* in (i).

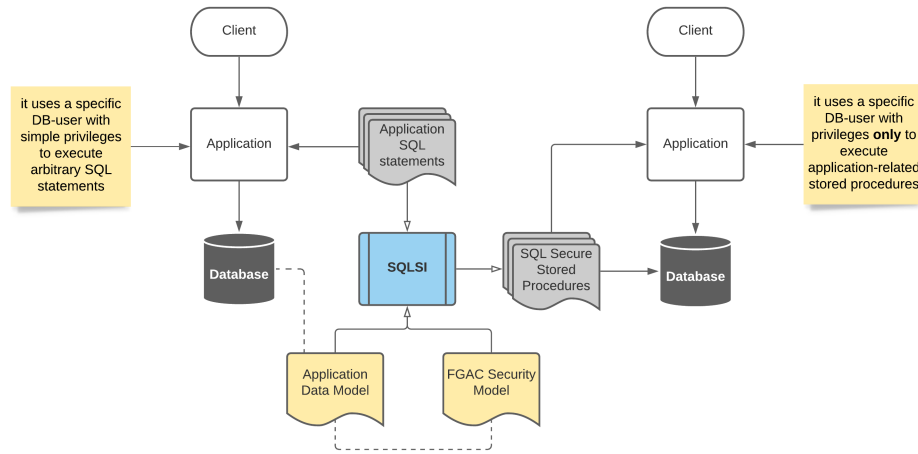


Fig. 1: Traditional approach vs. SQLSI approach

2.2 SQLSI realization vs low-code realization

In this paper, we focus on realizing the first functionality of SQLSI, in a more low-code fashion. To distinguish this with the former realization, let us call our new approach $\overline{\text{SQLSI}}$. Then, the following work of $\overline{\text{SQLSI}}$ will be primary:

- (i) Define the metamodel for data-model, security-model, and their domain-specific language.
- (ii) Define customized-model of relational database schema and the transformation from data-model to this relational database schema model.
- (iii) Define the code-generation from my relational database schema model to actual SQL database script.
- (iv) Define the transformation from the *user-input* version to the *normalized* version of the security-model.
- (v) Define the code-generation from security-model to a set of SQL authentication functions.

Figure 2 shows the workflow in a nutshell, including the technology used, whose description will be explained in Section 3 and Section 4.

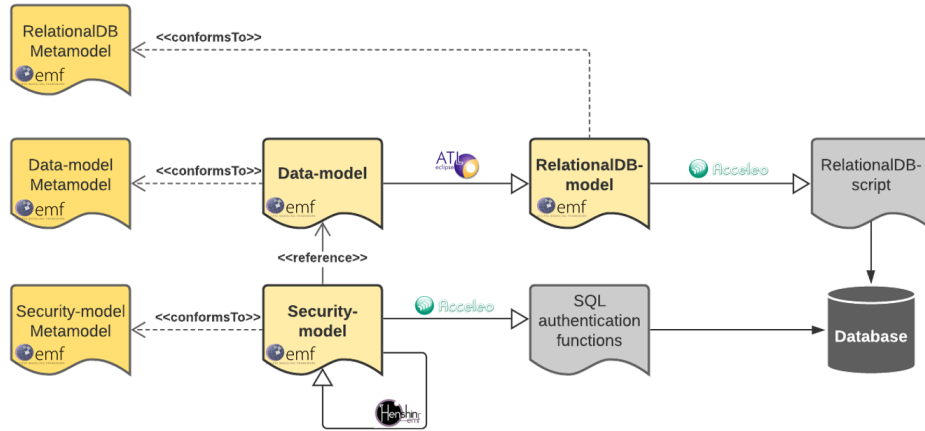


Fig. 2: Proposed approach - $\overline{\text{SQLSI}}$ in a nutshell

2.3 Related work

□ presents a Java implementation of the whole SQLSI component, including the SQL authorization functions generation. In particular, the datamodel and securitymodel are represented as Java object hierarchies and consequently, every model transformation and manipulation was implemented manually through the help of the visitor pattern. Indeed, this is a very ad-hoc approach and less modern, compared to the “one-click-of-a-button” implementation in this report.

[?,?] are the latest attempts to implement intelligent web-based editors for `datamodel` and `securitymodel`. In particular, these implementations are the final products of bachelor theses. On the one hand, these approaches are, again, very ad-hoc, since any change in the model requires (a handful of) manual changes in the source code. And on the other hand, due to the time allotted for the final thesis, the students do not have enough time to thoroughly understand the idea behind these models and therefore implemented in a way that is very hard to maintain or extend. In my approach, we avoid both drawbacks by using the available modelling technology, for instances, Xtext and ATL.

3 The SQLSI Metamodels

3.1 Input Metamodels

Metamodel for data-models For $\overline{\text{SQLSI}}$, a data-model contains entities and associations between them. An entity may have properties which are attributes or associations-ends.

The data-model metamodel for $\overline{\text{SQLSI}}$ is shown in Figure 3. The `DataModel` is the root element: it has a `name` and contains a set of `Entity`s. Every `Entity` represents an entity in the data model: it has a `name` and contains a set of `Property`(-ies). Every `Property` can be either an `Attribute` or an `AssociationEnd`.

- Each `Attribute` represents an attribute of the entity: its unique-ability depends on a boolean value `isUnique` and its `type` is either `String` or `Integer`.
- Each `AssociationEnd` represents an association between its container `Entity` with another `Entity`¹: its `Multiplicity` is either `MANY` or `ONE`. Finally, each `AssociationEnd` is linked to its opposite `AssociationEnd`, and also with its `targetEntity`.

Metamodel for security-models For $\overline{\text{SQLSI}}$, a security-model of a data-model contains a list of “fine-grained” rules which represent the authorization constraints for a data-model. Each rule formally states that under which *roles*, under which *actions*, under which *conditions*, the user can access which *resources*.

The security-model for $\overline{\text{SQLSI}}$ is shown in Figure 4. The `SecurityModel` is the root element, it has a `name` and contains a set of `ProtectedResources`, a set of `Rules` and a set of `Roles`. Naturally, every instance of `SecurityModel` is associated with an instance of `DataModel`.

- Each `ProtectedResource` represents a resource that the security-model intend to enforce fine-grained access control checks. Every protected resource has a `name` and is linked to a `Property` of the referenced `DataModel`.
- Each `Role` represents a role in the application: it has a `name`. Naturally, all the roles must referred to the *user-Entity* instance from the referenced `DataModel`.

¹ not necessary be a different one

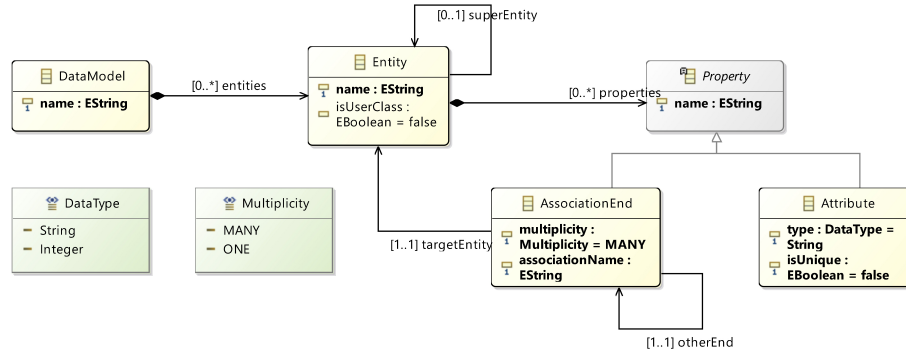


Fig. 3: SQLSI metamodel for data models.

- Each **Rule** represents a FGAC-rule in the application: it has a **name** and accommodates with an **action** (chosen from CRUD). Furthermore, every rule also refers to the resources that it intends to protect along a set of **Authorization** conditions, whose meaning is to state the scenarios when the *resource* is allowed to perform the *action*.²

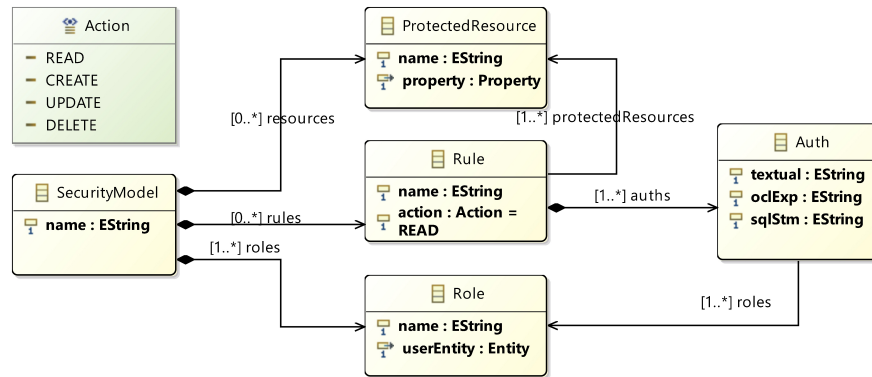


Fig. 4: SQLSI metamodel for security models.

² In this model, we allow user to write authorization constraints under three different means: either textual, OCL constraint or SQL boolean statement. Furthermore, each authorization constraint can be enforced for a set of **Roles**.

3.2 Output Metamodels

Metamodel for Relational Schema The relational schema metamodel is shown in Figure 5. It is self-explanatory, therefore its description shall be omitted here to save some spaces.

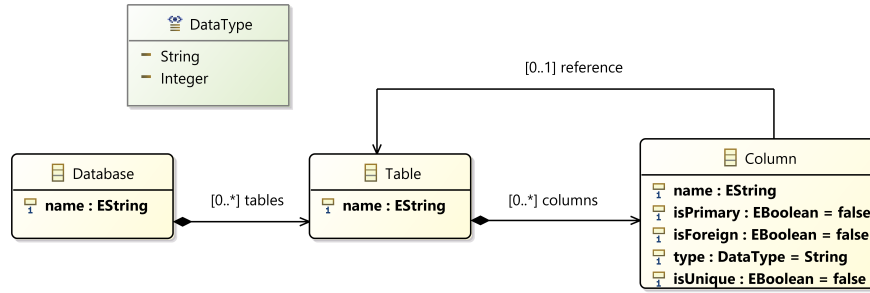


Fig. 5: $\overline{\text{SQLSI}}$ metamodel for relational database.

4 The $\overline{\text{SQLSI}}$ Language and Design

For the sake of clarity, let us continue this section of $\overline{\text{SQLSI}}$ language and design by associating it with a running example.

4.1 A Simple University Management System

Considering a simple UML class diagram **University** in Figure 6 containing four entities: **RegisterUser**, for representing register users; **Lecturer**, for representing lecturers; **Student**, for representing students; and **Course**, for representing the courses in the university. Every **RegisterUser** have a **name** and a unique **email**. A **Lecturer** is a **RegisterUser** with a **salary** attribute, for representing his/her monthly income. A **Student** is also a **RegisterUser** with a year **intake**. A **Course** has a **name** and a **year**. Every **Course** is taught by exactly one **Lecturer** and can have none or many **Students**. Additionally, every **Lecturer** can teach none or many **Courses** as well as every **Student** can enroll in none or many **Courses**.

4.2 Domain-Specific Language for the Input Models

DSL for data-model Figure 7 shows **University** written in our domain specific language. For the interested readers, the specification of this language can be seen at here.³ The additional features (e.g. validating, error-handling, scoping) will be leave for future works.

³ This language syntax is inspired by the ActionGUI[] datamodel language specification.

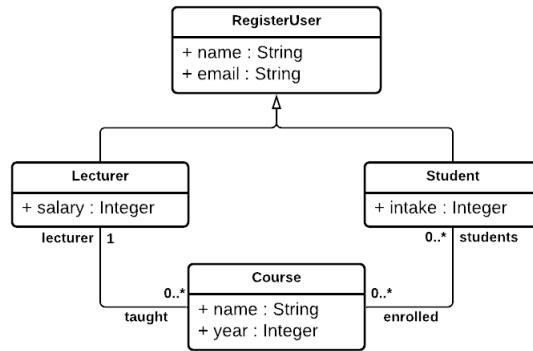


Fig. 6: University UML class diagram

```

DataModel University :
user entity RegUser {
  attribute name String,
  attribute email unique String
},
entity Student extends RegUser {
  attribute intake Integer,
  association Course[*] enrolled
  oppositeTo "Course.students" in Enrollment
},
entity Lecturer extends RegUser {
  attribute salary Integer,
  association Course[*] taught
  oppositeTo "Course.lecturer" in Teaching
},
entity Course {
  attribute name String,
  attribute year Integer,
  association Student[*] students
  oppositeTo "Student.enrolled" in Enrollment,
  association Lecturer[1] lecturer
  oppositeTo "Lecturer.taught" in Teaching
}

```

Fig. 7: University written in our data-model DSL

DSL for security-model Consider the following security-model for **UniversityDM**:

- Protected resources: **name** and **email** of the **RegisterUser**, **salary** of the **Lecturer** and **intake** year of the **Student**.⁴
- Role: There are three roles, namely, the role **Administrator**, **Lecturer** and **Student**.
- Permissions: each item represents a rule to enforce, in which is textually explained, follows by the OCL expression.
 - Any user can read any user's basic information. Formally,
`true`
 - Administrator can read any student intake year. Formally,
`true`
 - Student can read its classmates' intake year. Formally,
`caller.enrolledCourses→exists(c|c.students→includes(self))`⁵
 - Lecturer can read his\her students' intake year. Formally,
`caller.taughtCourses→exists(c|c.students→includes(self))`
 - Lecturer can read its own salary. Formally,
`caller = self`

Figure 8 shows a fragment of the above security-model written in our domain specific language. For the interested readers, the specification of this language can be seen at [here](#). The additional features (e.g. validating, error-handling, scoping) will be leave for future works.

4.3 Model Transformation: from Data Model to Relational Schema

Some transformation remarks Since transformation from UML class diagram to relational database is widely considered as a “Hello world!” example in the world of model transformation. The detail description of this section have been reduced to save some spaces for other interesting parts.

In short, our model transformation from **DataModel** to **RelationalDatabase** is implemented in ATLAS Transformation Language. It is very simple and indeed very similar to the example provided in the course. The main difference is that: during the transformation, it (i) creates a **Role-Table** and (ii) link this with the **Table** that correspond to the user-**Entity** via an association **Table**.

Code generation: Relational database to executable script In addition to the above transformation, we also provide a database schema generation using Aceleo. In this current version, we only support database schema generation that is executable on MySQL Server⁶ but it can be easily extended to other

⁴ Please note that the properties in **DataModel** which does not appear here simply mean that they are unprotected, i.e. they can be accessed by all means.

⁵ This is an OCL constraint in which **self** represents the user (object) whose information is accessed and **caller** represents the user (subject) whose are accessing the information.

⁶ Working version: 8.0.16 MySQL Community Server.


```

Rule readStudentSpecificInfo {
  action READ (studentIntake)
  auths {
    roles (Administrator)
    condition: {
      textual "Administrator can read
any student intake year"
      oclExp "true"
      sqlStm "TRUE"
    },
    roles (Student)
    condition: {
      textual "Student can read its classmates' intake year"
      oclExp "caller.enrolled->exists(c|c.students->includes(self))"
      sqlStm "caller = self"
    },
    roles (Lecturer)
    condition: {
      textual "Lecturer can read its students' intake year"
      oclExp "caller.taught->exists(c|c.students->includes(self))"
      sqlStm "EXISTS (SELECT 1
FROM (SELECT * FROM teaching WHERE lecturer = caller) as TEMP1
JOIN (SELECT * FROM enrollment WHERE students = self) as TEMP2
ON TEMP1.taught = TEMP2.enrolled)"
    }
  }
},
}

```

Fig. 8: A snippet of University security-model

language by create different configuration files.⁷ For interested readers, here shows the generated University MySQL database schema script.

4.4 Model Manipulation: Security Model

Rationale From the user's perspective, it is most often easier to create an FGAC rule of structure: “*role r* is able to perform *action a* on *resource p* under *condition c*”. In contrast, from the system's perspective, it is easier to enforce the FGAC policy if all the rules follow this order: *action-resource-role-condition*. For this reason, a model manipulation for security-model is needed. Since refining is not well-supported in ATLAS Transformation Language, Henshin has been chose for this task.

Remarks In short, the manipulation procedure is as follows:

1. For every authentication condition that applies for two different roles, we split those roles by create a copy of the original authentication condition. This will be applied exhaustively.

⁷ Since the relational database meta-model respects the standard of [], we should be able to provide generation for any given relational database management systems (i.e. Oracle, PostgreSQL, etc.), as long as it follows the same standard.

2. Next, for every rule that contains more than one authentication condition, we split those conditions by create a copy of the original rule. This will be applied exhaustively.
3. Then, for every rule that protects more than one resource, we split those protection by create a copy of the original rule. This will be applied exhaustively.
4. Finally, for every rule that protects the same resource and have the condition that applies for the same role, we merge them together.

Due to the space limit, the main unit is displayed in Figure 9 . For interested readers, the full transformation can be visited at .



Fig. 9: The main unit in our transformation

Code generation from Security Model to a Secure Authorization SQL-function In addition to the above transformation, we also provide a SQL authorization functions code generation using Acceleo. In this current version, we only support SQL code generation that is executable on MySQL Server⁸For interested readers, here shows the generated **University** SQL authorization functions artifact.

5 Conclusions and Future Work

⁸ Working version: 8.0.16 MySQL Community Server.